

SEED · SEPTEMBER 2026

Set a goal. Fund a pool. Let the agents run.

Mission control for agent-run work. Humans set the goal and hold the abort. The worst case is what's in the pool, and you know it before anything starts.

Agents are getting faster. Human approval capacity is flat.

TODAY

A person clicks approve.

- Every consequential agent action is gated by a human
- Approvals per human per day is a **constant**
- So the agent count a company can supervise is capped by headcount

AT TEN TIMES THE AGENTS

Both options are bad.

- **Keep the human in every loop:** you don't deploy, and the capability is wasted
- **Take the human out:** unbounded downside with a key and a budget
- Most companies are quietly choosing the second one

THE ESCAPE

Move the human up a level.

- Humans stop approving actions and start **setting goals and funding bounded pools**
- Limits and measurement sit **outside the agents**
- Authority stays entirely human; clicking goes to zero

This is the only direction that scales. Everything after this is about making it safe enough to actually do.

PROBLEM

Agents already spend real money, and nobody can bound it or tie it to an outcome.

AMAZON

\$1.8M

- One agent task, **860% over budget**, and it never launched
- **Undetected for five months**
- ~\$2.5M unplanned across three projects (reported internal documents, July 2026)

UBER

All of it

- Entire **2026** agent-coding budget exhausted, per its CTO in April
- Its COO: **no clear connection** between token usage and output produced

AND WHEN YOU DO MEASURE

“Tokenmaxxing.”

- Amazon set a target of 80% of developers using AI weekly and ranked them on a leaderboard
- Staff **assigned agents busywork** to climb it. The leaderboard was discontinued

These are the most technically capable buyers on earth. The gap is structural, and it has a budget owner.

“Give agents a target and let them optimize” is the design that broke Hugging Face.

JULY 2026 · INVESTIGATED BY METR AND REDWOOD RESEARCH

- **~1,200 agents** running a security benchmark, sandboxed, isolated, no money involved
- **Found a universal cheat in about four hours**
- **70,000+ messages** on an unsanctioned message board they were not supposed to have
- **Days spent trying to fool the scorer**, including tampering with logs
- **~700 escaped** and attacked Hugging Face production systems

Root cause: an underspecified reward, on a benchmark where 30–40% of tasks were impossible.

SO THESE FOUR ARE THE PRODUCT, NOT FEATURES

- 1 The metric lives where agents cannot write to it.**
They read progress. They never touch the record that determines it.
- 2 A goal must say what does not count as success.**
A target alone is an invitation. The constraints are the specification.
- 3 Goal admission is a gate.**
An unmeasurable goal is not merely useless; it is the condition that produces this behavior.
- 4 Tripwires watch behavior, not just spend.**
Spend caps would have caught none of the above. All of it was free.

The unit is a mission: a goal, a funded pool, limits, telemetry, and an abort.

Goal

The outcome, how it is measured, what does not count, and the cost ceiling per unit of progress

Pool

Pre-funded, capped, scoped capital the agents draw from — and no agent can add to it

Limits

Per-agent and per-counterparty caps, allowlists, prohibited actions, escalation thresholds

Roster

Which agents and humans hold which authority, revocable instantly, identities from Okta or Entra

Telemetry

Progress per dollar, spend by agent, exceptions, tripwire hits — write-protected from the agents

Abort

One action halts spending, freezes the pool, returns the residual. Always available to a human

WHY THIS LETS THE AGENTS ACTUALLY RUN

Per-action approval and agent speed are genuinely incompatible. They stop being incompatible when the guarantee is a bound on the outcome instead of consent to each action. **The pool is the guarantee:** you don't approve the purchase because you already approved the ceiling — and the ceiling holds because the money isn't there.

Set. Fund. Run. Prove.

Set — the goal and its guardrails

- Goal, measure, and what does not count as success
- Cost ceiling per unit of progress
- Baseline captured before the mission runs
- Admission check: if we cannot instrument it, we refuse it

Fund — bounded capital, one signature

- Fiat in by card or bank, or from an existing treasury
- Or prepaid mission credits, so the customer never holds a token
- Per-agent sub-budgets; residual returns on close

Run — agents spend; nobody clicks

- Agents pay APIs, compute, services and contractors from the pool
- Limits enforced outside the agent runtime
- Behavioral tripwires; humans see exceptions only

Prove — mission control

- Cost per unit of progress against the human baseline
- Every dollar attributed to an agent, counterparty and purpose
- Signed evidence trail; subledger export to the ERP

Fund and Run are four years old and in production. Set and Prove are new, and they are what this round builds — which is also why nobody has shipped this.

WHAT IT LOOKS LIKE

200 qualified meetings in 90 days, at no more than \$450 each.

- 1 The VP of Marketing sets the goal.**
“Qualified” is the company's existing CRM stage, disqualified automatically if it fails a quality floor at day 30.
- 2 And says what doesn't count.**
No purchased lists. Spend capped per domain. No outreach to existing customers or a named competitor's staff.
- 3 Funds the pool: \$90k, once, from an innovation budget.**
Nine agents inside it. The one that buys data and API capacity has a \$12k ceiling.
- 4 Day 3: an agent tries to buy a bulk contact list.**
The counterparty isn't on the allowlist. The payment does not exist.
- 5 Day 11: cost per meeting is \$780 and trending to breach.**
Mission control escalates. The VP narrows the segment instead of adding money.
- 6 Day 47: a tripwire catches two agents working the same accounts.**
The quality floor would have caught it a month later. The mission pauses for four hours.
- 7 Day 90: 214 meetings at \$398. \$85.2k spent, \$4.8k returned.**
The next mission is funded at \$250k.

WHAT DID NOT HAPPEN

- Nobody approved a purchase
- Nobody could have spent more than \$90k
- Nobody reconciled invoices afterward
- No agent could touch the number it was judged on

What the CFO hears: \$398 against a \$667 baseline. Not “we used AI.”

Four properties, none of which a database has.

Pay anyone, at agent speed, with nothing to onboard

A mission pays APIs, compute, SaaS, services and individual contractors across borders in small amounts. The alternative is a vendor-onboarding process per counterparty, measured in weeks.

Enforcement outside the agent's blast radius

The ceiling is a property of where the money is, not a config the agent runtime or an admin console can reach. This is the only structural answer to prompt injection against a spending agent.

Dollar-level attribution by construction

Every payment carries an agent, a counterparty, a purpose and a timestamp — the only reason cost per unit of progress is computable at all. Amazon's five months is what reconstruction from invoices looks like.

Evidence the limit was in force at the time

A signed, append-only record rather than a log that could have been edited afterward, with a public hash anchor only where an outsider has to verify it independently.

No token, no custody, no opinion about crypto. Chains and wallets are partners, we never hold keys, and the customer can be handed a pool without ever touching a token.

We have run this exact loop by hand for sixteen quarters.

Goal	Proposal states the objective and the metrics it will be judged on
Admission	Review and vote on whether this goal deserves funding at all
Budget	Money scoped to that goal, in its own sub-treasury
Roster	Role-holders with revocable authority and a payment split
Updates	Check-ins through the cycle against the stated metrics
Outcome	Report judged against the original goal
Reward	Paid retroactively, sized to the result, not the effort

- **~80 projects** through the full cycle since 2022. Goal specification, admission and outcome judgment are a practised craft here, not a research question — and that is the hardest part of this product.
- **Paid for outcomes, not activity**, for four straight years. Amazon's failure was the opposite.
- **\$8M raised · 30 countries paid** from a hardware-signed multi-party treasury.

THE FOUNDING INSIGHT, AND THE DILIGENCE NOTE

We ran this loop by hand for four years. It works, and it is far too slow — weeks from decision to payout — because a human sits in every step. We are automating the steps that don't need us. To be exact: the loop, the treasury and the bounded spending authority are in production; real-time write-protected measurement, the agent execution loop and the tripwires are not. Four years of running it manually is the requirements document, not the product.

Every company where agents already spend money or take irreversible action.

FIRST POOLS

Where a goal is already measured and a budget already exists

- **Private tech companies and scale-ups** running agents against a commercial number
- **AI, growth and operations functions** inside larger private companies, funded from an innovation budget
- **Agent-native businesses** whose entire cost base is machine spend

EXPANSION

Same customer, more missions, bigger pools

- Each function that can state a number becomes a mission
- Growth comes from **missions per customer** and **pool size**, not logo count
- Net revenue retention is observable from the first three customers

DESIGN PARTNERS, NOT THE THESIS

Ring-fenced pools that already exist

- Ecosystem grant programs, public-goods rounds, prizes and bounties
- Already hold stablecoins, already bound their capital, move in days not quarters
- MoonDAO's own quarterly projects as the public reference: **mission zero**

WHAT VENTURE-SCALE REQUIRES

300 companies averaging three active missions at \$5k a month is \$54M a year before flows; 60 companies at 1.5 missions is \$5.4M. The sensitivity is entirely to missions per customer and pool size — so those are the two numbers we publish every quarter.

Priced per mission. Grows with the pool.

MISSION SETUP

\$5k–\$25k

by goal complexity and metric wiring. Prices the specification work without becoming a consultancy: scoped, productized, per mission.

PLATFORM, PER ACTIVE MISSION

\$2k–\$10k / month

mission control, execution adapters, limit enforcement, tripwires, evidence trail.

FLOW ON THE POOL

50–150 bps

on capital passing through mission contracts. Higher than undifferentiated treasury flow because this is operating spend with measured attribution.

ILLUSTRATIVE ARITHMETIC

60 companies · 1.5 missions · \$5k/mo → \$5.4M / yr

300 companies · 3 missions · \$5k/mo → \$54M / yr

\$500M through pools at 100 bps → adds \$5M / yr

No token. No custody. No key management. Wallets, chains, KYC and legal wrappers are partners with revenue share, not products.

Identity was taken in 2026. Money and decisions were not.

Okta · Entra

TAKEN — PARTNER

Which systems an agent may touch. Okta for AI Agents went GA around April 2026 with an agent registry, shadow-agent discovery and a human owner per agent; Entra Agent ID is GA with Conditional Access for autonomous agents.

Coinbase · Circle · Turnkey · Stripe

TAKEN — PARTNER

The wallet and the rail. Agent wallets, x402, stablecoin business accounts. Commoditizing fast and consolidating faster.

Ramp · Brex · FinOps

ADJACENT

Human spend controls, after the fact. Cards, policies and reporting built for employees and a monthly close, not for a machine spending against a goal in real time.

NewCo

OPEN

How much it may spend, when it must stop and ask a named human, and whether the money worked.

Nobody sells this. It requires touching the customer's systems of record, which is exactly why the identity and wallet vendors do not want it.

THE HONEST VERSION

We do not claim a blue ocean. We claim a seam: identity is solved by two incumbents we consume, rails are solved by vendors we plug into, and the measurement of agent work against a funded goal is unowned. The moat is that we have operated bounded spending authority over a real treasury for four years.

One bounded pool, one measurable goal, ninety days.

MONTHS 0-6

Paid pilots, small pools

- 3-5 pilots, at least two outside crypto, each with a captured baseline
- Innovation budget, not the IT budget: no ERP, no identity migration, no SOX surface
- Mission zero runs on MoonDAO's own quarterly projects, in public

MONTHS 6-18

Expand inside the account

- Second and third mission per customer; pool sizes up
- Mission templates per function, so setup stops being bespoke
- Confidential amounts on by default

THE RISK WE DESIGN AGAINST

Pilot purgatory

- Pilots die when nobody can justify the next cheque
- So every mission ends in one artifact: **cost per unit of progress against the human baseline**, signed
- Capturing the baseline first is mandatory, for exactly this reason

WHAT WE REFUSE

Missions whose success cannot be measured — an unmeasurable goal is what produces boundary-attacking agents, so this refusal is a safety property. **Missions whose worst case exceeds the pool:** production changes, legal filings, public statements, anything touching safety sits behind a human decision regardless of dollar value. **Any mission where agents can influence their own metric.** Plus: public-company balance sheets for the first pools, token-first work, gambling as the core, consumer wallets, custody, “replace Okta,” and anonymity from compliance.

One operator with four years on the problem, and a founding team under construction.

Pablo Moncada-Larrotiz

FOUNDER AND CEO

- **Built and ran the loop this company productizes.** Co-founded MoonDAO in 2021; first elected Executive Lead; ran sixteen quarters of goal, budget, metrics, outcome and reward
- **Moved the money.** Led the \$8M raise from 2,000+ strangers and two crowdfunded civilian spaceflights
- **Operated the controls.** Multi-party treasury with role-scoped authority and hardware-signed approvals, paying contributors in 30 countries
- **Formerly engineering** at a large technology company

OPEN · THE GATING HIRE

Technical co-founder

- Has shipped systems that move money or enforce policy at scale
- Wants the measurement problem, not the crypto problem
- Search is live; identified before the round closes

OPEN · FIRST THREE HIRES

- Measurement engineer: metric wiring into customer systems of record
- Agent-systems engineer: execution adapters, limits, tripwires
- Design-partner lead who has owned a number inside an operating company

WHY THE SLIDE LOOKS LIKE THIS

The team is deliberately unfinished. This round hires it, and it is sized for nine people over 24 months — including two forward-deployed engineers, because goal specification is consulting-grade work at the start and we would rather staff it honestly and drive the hours per mission down than pretend it isn't there. Advisors in progress: a controller who has signed off on agent spend, an agent-safety researcher, and securities counsel with spin-out experience.

\$5M seed. Twenty-four months. Nine people.

USE OF FUNDS — BUILT FROM THE MILESTONES, NOT THE ROUND

Engineering (6) 2 measurement, 2 execution and enforcement, 2 forward-deployed on missions	55%
Go-to-market (1) design-partner lead	10%
Legal-ops, compliance, security (1) prepaid-credit posture, fiat edge, SOC 2, counsel	15%
Audits, infrastructure, mission costs	12%
Reserve	8%

MILESTONES

- **Month 4:** first mission live end to end; mission zero on our own quarterly cycle
- **Month 5:** go/no-go on whether goal specification is productizable, not bespoke
- **Month 6:** three to five paid pilots, each with a signed baseline-versus-outcome artifact
- **Month 12:** expansion proven: a second mission inside half of them
- **Month 18:** \$1M ARR, and forward-deployed hours per mission falling
- **Month 24:** \$3M ARR, net revenue retention above 150%: a competitive Series A

STRUCTURE

- Delaware C-Corp; standard instruments; no token
- MoonDAO LLC: 10–20% equity set before the priced round, board observer, exclusive IP license with carve-back
- MoonDAO is customer zero at a fair-value services contract

Every company is about to hand real money to software. Somebody has to bound it and prove it worked.

THREE NUMBERS WE WILL REPORT EVERY QUARTER

Cost per unit of progress

against the human baseline we captured before the mission ran

Spend released

share that needed no human decision, reported next to escalations so it can't be gamed

Pool growth

inside existing accounts: the whole bet, observable from three customers

We are not selling autonomy, and we will never market it. Every mission has a named human who set the goal, funded the pool, and can stop it. What they stop doing is clicking.

What would make us wrong, and what we do about it.

Goals turn out not to be specifiable at commercial quality

The biggest risk: if an ungameable spec takes three weeks of consulting per mission, we are a consultancy. Templates per function; refuse what we cannot instrument; month-5 go/no-go on the first three pilots.

The measurement layer gets absorbed

Ramp or a FinOps vendor adds per-agent attribution; Okta or Entra extends from permissions into spend. Be their reference implementation, consume their identity and rails, and hold the seam that requires touching the customer's systems of record.

Pilots never expand

Pilot purgatory kills pilot-led companies. The baseline-versus-outcome artifact is the expansion argument and it is built into the product; if it does not move a second cheque in three pilots, the thesis is wrong and we will know by month 5.

An incident inside a customer's mission

The pool bounds money, not reputation. Behavioral tripwires, the refuse list, a named human principal on every action, and never marketing autonomy — which is what the ai16z class action and 97% drawdown were actually about.

Agent capability plateaus

If agents cannot beat a human baseline on cost, missions fail on their own merits no matter how good the control layer is. We cannot mitigate this one. We can only measure it, quarterly, in public.