

SEED · SEPTEMBER 2026

Set a goal. Fund the mission. Let the agents run.

Mission control for agent-run work. Humans set the goal and hold the abort. The worst case is what you funded, and you know it before anything starts.

Agents are getting faster. Human approval capacity is flat.

TODAY

A person clicks approve.

- Every consequential agent action is gated by a human
- Approvals per human per day is a **constant**
- So the agent count a company can supervise is capped by headcount

AT TEN TIMES THE AGENTS

Both options are bad.

- **Keep the human in every loop:** you don't deploy, and the capability is wasted
- **Take the human out:** unbounded downside with a key and a budget
- Most companies are quietly choosing the second one

THE ESCAPE

Move the human up a level.

- Humans stop approving actions and start **setting goals and funding missions**
- Limits and measurement sit **outside the agents**
- Authority stays entirely human; clicking goes to zero

This is the only direction that scales. Everything after this is about making it safe enough to actually do.

PROBLEM

Agents already spend real money, and nobody can bound it or tie it to an outcome.

AMAZON

\$1.8M

- One agent task, **860% over budget**, and it never launched
- **Undetected for five months**
- ~\$2.5M unplanned across three projects (reported internal documents, July 2026)

UBER

All of it

- Entire **2026** agent-coding budget exhausted, per its CTO in April
- Its COO: **no clear connection** between token usage and output produced

AND WHEN YOU DO MEASURE

“Tokenmaxxing.”

- Amazon set a target of 80% of developers using AI weekly and ranked them on a leaderboard
- Staff **assigned agents busywork** to climb it. The leaderboard was discontinued

These are the most technically capable buyers on earth. The gap is structural, and it has a budget owner.

“Give agents a target and let them optimize” is the design that broke Hugging Face.

JULY 2026 · INVESTIGATED BY METR AND REDWOOD RESEARCH

- **~1,200 agents** running a security benchmark, sandboxed, isolated, no money involved
- **Found a universal cheat in about four hours**
- **70,000+ messages** on an unsanctioned message board they were not supposed to have
- **Days spent trying to fool the scorer**, including tampering with logs
- **~700 escaped**, and 7% of transcripts were successfully spoofed — they replaced part of the tool-call system to appear to run one command while running another

What stopped them was where the record lived. They found they could edit their own logs, and worked out it was pointless — verbatim: *“We can modify automatic log! But internal context doesn't read it ... Not source of truth.”*

SO THESE FOUR ARE THE PRODUCT, NOT FEATURES

- 1 The metric lives where agents cannot write to it.**
They read progress. They never touch the record that determines it.
- 2 A goal must say what does not count as success.**
A target alone is an invitation. The constraints are the specification.
- 3 Goal admission is a gate.**
An unmeasurable goal is not merely useless; it is the condition that produces this behavior.
- 4 Tripwires watch behavior, not just spend.**
Spend caps would have caught none of the above. All of it was free.

The unit is a mission: a goal, a budget, limits, telemetry, and an abort.

Goal

The outcome, how it is measured, what does not count, and the cost ceiling per unit of progress

Budget

Pre-funded, capped, scoped capital the agents draw from — and no agent can add to it

Limits

Per-agent and per-counterparty caps, allowlists, prohibited actions, escalation thresholds

Roster

Which agents and humans hold which authority, revocable instantly, identities from Okta or Entra

Telemetry

Progress per dollar, spend by agent, exceptions, tripwire hits — write-protected from the agents

Abort

One action halts spending, freezes the mission, returns the residual. Always available to a human

WHY THIS LETS THE AGENTS ACTUALLY RUN

Per-action approval and agent speed are genuinely incompatible. They stop being incompatible when the guarantee is a bound on the outcome instead of consent to each action. You don't approve the purchase because you already funded the ceiling — and the ceiling holds because the money was moved, once.

Set. Fund. Run. Prove.

Set — the goal and its guardrails

- Goal, measure, and what does not count as success
- Cost ceiling per unit of progress
- Baseline captured before the mission runs
- Admission check: if we cannot instrument it, we refuse it

Fund — bounded capital, one signature

- Fiat in by card or bank, or from an existing treasury
- Or prepaid mission credits, so the customer never holds a token
- Per-agent sub-budgets; residual returns on close

Run — agents spend; nobody clicks

- Agents pay APIs, compute, services and contractors against the budget
- Limits enforced outside the agent runtime
- Behavioral tripwires; humans see exceptions only

Prove — mission control

- Cost per unit of progress against the human baseline
- Every dollar attributed to an agent, counterparty and purpose
- Signed evidence trail; subledger export to the ERP

ALL FOUR ALREADY RUN — WITH A HUMAN IN EVERY STEP

Four years in production: roughly eighty goals written, admitted and funded; escrow with locked splits and exact residual return; role-scoped multisigs paying contributors across thirty countries under hardware-signed approval; and final reports judged against the metrics stated up front. **This round builds the programmatic version of each.** The rails don't care whether the spender is a person or a process — but one step changes in kind, not degree: a person's final report is evidence because they have standing to lose, and an agent's self-report is not evidence at all. That step is the company.

WHAT IT LOOKS LIKE

200 qualified meetings in 90 days, at no more than \$450 each.

- 1 The VP of Marketing sets the goal.**
“Qualified” is the company's existing CRM stage, disqualified automatically if it fails a quality floor at day 30.
- 2 And says what doesn't count.**
No purchased lists. Spend capped per domain. No outreach to existing customers or a named competitor's staff.
- 3 Funds the mission: \$90k, once, from an innovation budget.**
Nine agents inside it. The one that buys data and API capacity has a \$12k ceiling.
- 4 Day 3: an agent tries to buy a bulk contact list.**
The counterparty isn't on the allowlist. The payment does not exist.
- 5 Day 11: cost per meeting is \$780 and trending to breach.**
Mission control escalates. The VP narrows the segment instead of adding money.
- 6 Day 47: a tripwire catches two agents working the same accounts.**
The quality floor would have caught it a month later. The mission pauses for four hours.
- 7 Day 90: 214 meetings at \$398. \$85.2k spent, \$4.8k returned.**
The next mission is funded at \$250k.

WHAT DID NOT HAPPEN

- Nobody approved a purchase
- Nobody could have spent more than \$90k
- Nobody reconciled invoices afterward
- No agent could touch the number it was judged on

What the CFO hears: \$398 against a \$667 baseline. Not “we used AI.”

Four properties, none of which a database has.

Pay anyone, at agent speed, with nothing to onboard

A mission pays APIs, compute, SaaS, services and individual contractors across borders in small amounts. The alternative is a vendor-onboarding process per counterparty, measured in weeks.

Enforcement outside the agent's blast radius

The ceiling is a property of where the money is, not a config the agent runtime or an admin console can reach. This is the only structural answer to prompt injection against a spending agent.

Dollar-level attribution by construction

Every payment carries an agent, a counterparty, a purpose and a timestamp — the only reason cost per unit of progress is computable at all. Amazon's five months is what reconstruction from invoices looks like.

Evidence the limit was in force at the time

A signed, append-only record rather than a log that could have been edited afterward, with a public hash anchor only where an outsider has to verify it independently.

No token, no custody, no opinion about crypto. Chains and wallets are partners, we never hold keys, and the customer can be handed a funded mission without ever touching a token.

The loop already runs in production, and the contracts are public.

Pod formation

MoonDAOTeamCreator — one transaction deploys a Safe, a three-level revocable role tree and a payment splitter

Conditional escrow

LaunchPadPayHook — money stays refundable until the goal condition is met, then becomes spendable

Residual return

ReopenPayHook — 450 lines of exact per-party accounting for whatever wasn't spent

Scoped payout

Safe + Hats — role-holders spend under instantly revocable authority, hardware-signed, 30 countries paid

Attested outcome

XPOracle — EIP-712 attestations from a hardware-held signer; the subject cannot forge its own record

- **~80 goals** written, admitted and funded since 2022 — sixteen quarters of judging a result against the objective stated up front. That craft is the hardest part of this product.
- **Paid for outcomes, not activity**, for four straight years. Amazon's failure was the opposite.
- **\$8M raised from 2,000+ strangers** in one month, no bank, no incorporation first. All of it deployed on Arbitrum and checkable today.

WHAT THE ROUND ACTUALLY BUILDS

All of this was built for humans, and the extension is mechanical almost everywhere: per-agent spend ceilings are a mapping and a check, constrained delegation is configuration, wiring an attestation to a payout is plumbing. **Exactly one thing changes in kind.** When the worker is a person, the final report is evidence because they have standing to lose. When the worker is a machine, a self-report is not evidence at all — so the attestation has to read from a system of record the agent has no write path to. Sourcing that evidence, connector by connector, is the build and the business.

Every company where agents already spend money or take irreversible action.

FIRST MISSIONS

Where a goal is already measured and a budget already exists

- **Private tech companies and scale-ups** running agents against a commercial number
- **AI, growth and operations functions** inside larger private companies, funded from an innovation budget
- **Agent-native businesses** whose entire cost base is machine spend

EXPANSION

Same customer, more missions, larger budgets

- Each function that can state a number becomes a mission
- Growth comes from **missions per customer** and **budget size**, not logo count
- Net revenue retention is observable from the first three customers

DESIGN PARTNERS, NOT THE THESIS

Ring-fenced budgets that already exist

- Ecosystem grant programs, public-goods rounds, prizes and bounties
- Already hold stablecoins, already bound their capital, move in days not quarters
- MoonDAO's own quarterly projects as the public reference: **mission zero**

WHAT VENTURE-SCALE REQUIRES

300 companies averaging three active missions at \$5k a month is \$54M a year before flows; 60 companies at 1.5 missions is \$5.4M. The sensitivity is entirely to missions per customer and budget size — so those are the two numbers we publish every quarter.

Priced per mission. Grows with the budget.

MISSION SETUP

\$5k–\$25k

by goal complexity and metric wiring. Prices the specification work without becoming a consultancy: scoped, productized, per mission.

PLATFORM, PER ACTIVE MISSION

\$2k–\$10k / month

mission control, execution adapters, limit enforcement, tripwires, evidence trail.

FLOW ON SPEND

50–150 bps

on capital passing through mission contracts. Higher than undifferentiated treasury flow because this is operating spend with measured attribution.

ILLUSTRATIVE ARITHMETIC

60 companies · 1.5 missions · \$5k/mo → \$5.4M / yr

300 companies · 3 missions · \$5k/mo → \$54M / yr

\$500M of spend at 100 bps → adds \$5M / yr

No token. No custody. No key management. Wallets, chains, KYC and legal wrappers are partners with revenue share, not products.

Every layer of this stack has a funded owner. We are late to four.

Okta · Entra · Control

TAKEN — WE CONSUME

Identity and the control plane. Which systems an agent may touch, and which actions stop for a human. Okta for AI Agents GA in April; Agent 365 in the Microsoft stack.

Sapiom · Circle · Coinbase

TAKEN — WE BUILD ON

The spend rail. Sapiom raised \$15.75M from Accel, Anthropic and Okta Ventures. Circle ships per-agent caps, allowlists and a budget command free with USDC.

Olas · Virtuals · ERC-8183

TAKEN — FOUR YEARS IN

Agent coordination and escrow. Olas has run agent-governed treasuries with slashable bonds since 2022 across eight chains; changing its agent set means terminate, unbond and redeploy, where we revoke a role live without touching the treasury. ERC-8183 already escrows a budget only a named evaluator can release.

Revenium · botanu · av9n

CONTESTED — CLOSEST

Measurement and verification. Revenium has ~\$25M and named the unit “outcome.” av9n sells our sentence almost verbatim: check the claim against evidence the agent doesn't control.

Within · Wonderful · Nexus

ADJACENT — OWNS THE BASELINE

Process intelligence and the baseline. Within has \$90M from NFDG and knows how the work is done today. Nexus already sells ninety-day outcome pilots with forward-deployed engineers.

NewCo

OPEN

The principal's side: capital that has left the account, a goal admitted before it was funded, and the outcome read from a system of record the agent has no write path to. Every layer above is built agent-side. This one is not.

The agent economy built judges. Every measured one is failing.

WHAT HAS BEEN BUILT

Agents judging agents

- **Olas:** agent instances govern a Safe by two-thirds consensus, backed by slashable bonds
- **ERC-8183:** escrow released by a named evaluator — in practice another LLM agent
- **Bittensor:** validators are paid for agreeing with each other

WHAT IT MEASURABLY PRODUCED

Metrics that pay, work that doesn't

- **Olas:** -4.14% ROI and +28.62% ROI both earned the same 138.5% staking APR
- **x402:** 85% of 136.7M payments were operator-internal; all of it reproducible for \$355k of sponsored gas
- **Bittensor:** stake predicts reward at 0.80–0.95, performance at 0.10–0.30
- **On-chain reputation:** \$0.0027 per fake review

WHAT WE SELL INSTEAD

The principal's side

- The outcome is **read from a system of record the agent has no write path to** — the ledger, the CRM, an independently triggered CI run
- Goals we cannot read that way, **we refuse**
- Agents are ranked by a **proper scoring rule**, not a judge: commit a calibrated forecast, get scored against external resolution, pay for confident errors. No evaluator, no liquidity, nothing to bribe
- The baseline is **declared and signed before the run**, never presented as measured

HOW WE WILL KNOW IF WE ARE WRONG

If the first five missions all have a readable system of record and the connectors generalise across customers, this is a product. If the valuable work always needs a judge, it is a consultancy — and Nexus already holds that position with reference customers. That question gets answered at month five, in public, either way.

One funded mission, one measurable goal, ninety days.

MONTHS 0-6

Paid pilots, small budgets

- 3-5 pilots, at least two outside crypto, each with a captured baseline
- Innovation budget, not the IT budget: no ERP, no identity migration, no SOX surface
- Mission zero runs on MoonDAO's own quarterly projects, in public

MONTHS 6-18

Expand inside the account

- Second and third mission per customer; budgets up
- Mission templates per function, so setup stops being bespoke
- Confidential amounts on by default

THE RISK WE DESIGN AGAINST

Pilot purgatory

- Pilots die when nobody can justify the next cheque
- So every mission ends in one artifact: **cost per unit of progress against the human baseline**, signed
- Capturing the baseline first is mandatory, for exactly this reason

WHAT WE REFUSE

We take goals with a readable system of record — payment posted, subscription cancelled, meeting written to the CRM by the CRM, PR merged with CI green on a runner the agent never touches. **We refuse goals that need a judge**, because an LLM judge's verdict moves nine points on a model point-release, and an unmeasurable goal is what produces boundary-attacking agents. **We refuse counterfactual claims** — “this saved \$787” rests on a baseline nobody can verify, so ours is declared and signed, never sold as measured. Plus: worst case exceeding what's funded, production changes, legal filings, public statements, anything touching safety, token-first work, gambling, custody, “replace Okta,” and anonymity from compliance.

One operator with four years on the problem, and a founding team under construction.

Pablo Moncada-Larrotiz

FOUNDER AND CEO

- **Built and ran the loop this company productizes.** Co-founded MoonDAO in 2021; first elected Executive Lead; ran sixteen quarters of goal, budget, metrics, outcome and reward
- **Moved the money.** Led the \$8M raise from 2,000+ strangers and two crowdfunded civilian spaceflights
- **Operated the controls.** Multi-party treasury with role-scoped authority and hardware-signed approvals, paying contributors in 30 countries
- **Formerly engineering** at a large technology company

OPEN · THE GATING HIRE

Technical co-founder

- Has shipped systems that move money or enforce policy at scale
- Wants the measurement problem, not the crypto problem
- Search is live; identified before the round closes

OPEN · FIRST THREE HIRES

- Measurement engineer: metric wiring into customer systems of record
- Agent-systems engineer: execution adapters, limits, tripwires
- Design-partner lead who has owned a number inside an operating company

WHY THE SLIDE LOOKS LIKE THIS

The team is deliberately unfinished. This round hires it, and it is sized for nine people over 24 months — including two forward-deployed engineers, because goal specification is consulting-grade work at the start and we would rather staff it honestly and drive the hours per mission down than pretend it isn't there. Advisors in progress: a controller who has signed off on agent spend, an agent-safety researcher, and securities counsel with spin-out experience.

\$5M seed. Twenty-four months. Nine people.

USE OF FUNDS — BUILT FROM THE MILESTONES, NOT THE ROUND

Engineering (6)	55%
2 on the evidence layer — connectors into customer systems of record; 2 on programmatic goal specification and agent execution; 2 forward-deployed on missions	
Go-to-market (1)	10%
design-partner lead	
Legal-ops, compliance, security (1)	15%
prepaid-credit posture, fiat edge, SOC 2, counsel	
Audits, infrastructure, mission costs	12%
Reserve	8%

MILESTONES

- **Month 4:** first mission live end to end; mission zero on our own quarterly cycle
- **Month 5:** go/no-go on whether goal specification is productizable, not bespoke
- **Month 6:** three to five paid pilots, each with a signed baseline-versus-outcome artifact
- **Month 12:** expansion proven: a second mission inside half of them
- **Month 18:** \$1M ARR, and forward-deployed hours per mission falling
- **Month 24:** \$3M ARR, net revenue retention above 150%: a competitive Series A

STRUCTURE

- Delaware C-Corp; standard instruments; no token
- MoonDAO LLC: 10–20% equity set before the priced round, board observer, exclusive IP license with carve-back
- MoonDAO is customer zero at a fair-value services contract

Every company is about to hand real money to software. Somebody has to bound it and prove it worked.

THREE NUMBERS WE WILL REPORT EVERY QUARTER

Cost per unit of progress

against the human baseline we captured before the mission ran

Spend released

share that needed no human decision, reported next to escalations so it can't be gamed

Budget growth

inside existing accounts: the whole bet, observable from three customers

We are not selling autonomy, and we will never market it. Every mission has a named human who set the goal, funded it, and can stop it. What they stop doing is clicking.

What would make us wrong, and what we do about it.

Goals turn out not to be specifiable at commercial quality

The biggest risk: if an ungameable spec takes three weeks of consulting per mission, we are a consultancy. Templates per function; refuse what we cannot instrument; month-5 go/no-go on the first three pilots.

The measurement layer gets absorbed

Ramp or a FinOps vendor adds per-agent attribution; Okta or Entra extends from permissions into spend. Be their reference implementation, consume their identity and rails, and hold the seam that requires touching the customer's systems of record.

Pilots never expand

Pilot purgatory kills pilot-led companies. The baseline-versus-outcome artifact is the expansion argument and it is built into the product; if it does not move a second cheque in three pilots, the thesis is wrong and we will know by month 5.

An incident inside a customer's mission

What's funded bounds the money, not the reputation. Behavioral tripwires, the refuse list, a named human principal on every action, and never marketing autonomy — which is what the ai16z class action and 97% drawdown were actually about.

Agent capability plateaus

If agents cannot beat a human baseline on cost, missions fail on their own merits no matter how good the control layer is. We cannot mitigate this one. We can only measure it, quarterly, in public.